OpenStreetMap Sample Project
Data Wrangling with MongoDB

Matthew Banbury

Map Area: Charlotte, NC, United States

https://www.openstreetmap.org/relation/177415
http://metro.teczno.com/#charlotte

1. Problems Encountered in the Map

After initially downloading a small sample size of the Charlotte area and running it against a provisional data.py file, I noticed three main problems with the data, which I will discuss in the following order:
- Over-abbreviated street names ("S Tryon St Ste 105")
- Inconsistent postal codes ("NC28226", "282260783", "28226")
- "Incorrect" postal codes (*Charlotte area zip codes all begin with "282" however a large portion of all documented zip codes were outside this region.*)

Over-abbreviated Street Names

Once the data was imported to MongoDB, some basic querying revealed street name abbreviations and postal code inconsistencies. I updated all substrings in problematic address strings, such that "S Tryon St Ste 105" becomes "South Tryon Street Suite 105".

Postal Codes

Postal code strings posed a different sort of problem, forcing a decision to strip all leading and trailing characters before and after the main 5-digit zip code. This effectually dropped all leading state

characters (as in "NC28226") and 4-digit zip code extensions following a hyphen ("282260783"). This 5-digit constriction benefits MongoDB aggregation calls on postal codes.

Regardless, after standardizing inconsistent postal codes, some altogether "incorrect" (or perhaps misplaced?) postal codes surfaced when grouped together with this aggregator:

Sort postcodes by count, descending

```
db.char.aggregate([{"$match":{"address.postcode":{"$exists":1}}},
{"$group":{"_id":"$address.postcode", "count":{"$sum":1}}},
{"$sort":{"count":1}}])
```

Here are the top two results, beginning with the highest count:

```
[ {"_id" : "29732", "count" : 103},
{"_id" : "28134", "count" : 27}, ...
```

Considering the relatively few documents that included postal codes, of those, it appears that out of the top ten, seven aren't even in Charlotte. That struck me as surprisingly high to be a blatant error, and found that the number one postal code and all others starting with "297" lie in Rock Hill, SC. So, I performed another aggregation to verify a certain suspicion...

Sort cities by count, descending

```
> db.char.aggregate([{"$match":{"address.city":{"$exists":1}}},
{"$group":{"_id":"$address.city", "count":{"$sum":1}}},
{"$sort":{"count":1}}])
```

And, the results, edited for readability:

```
[ {"_id" : "Rock Hill", "count" : 111},...
```

These results confirmed my suspicion that this metro extract would perhaps be more aptly named "Metrolina" or the "Charlotte Metropolitan Area" for its inclusion of surrounding cities in the sprawl. So, these postal codes aren't "incorrect," but simply unexpected.

2. Data Overview

This section contains basic statistics about the dataset and the MongoDB queries used to gather them.

File sizes

```
charlotte.osm ......... 294 MB
charlotte.osm.json .... 322 MB
```

Number of documents

```
> db.char.find().count()
1555851
```

Number of nodes

```
> db.char.find({"type":"node"}).count()
1471349
```

Number of ways

```
> db.char.find({"type":"way"}).count()
84502
```

Number of unique users

```
> db.char.distinct({"created.user"}).length
336
```

Top 1 contributing user

```
> db.char.aggregate([{"$group":{"_id":"$created.user", "count":{"$sum":1}}},
{"$sort":{"count":1}}, {"$limit":1}])
[ { "_id" : "jumbanho", "count" : 823324 } ]
```

Number of users appearing only once (having 1 post)

```
> db.char.aggregate([{"$group":{"_id":"$created.user", "count":{"$sum":1}}},
{"$group":{"_id":"$count", "num_users":{"$sum":1}}}, {"$sort":{"_id":1}},
{"$limit":1}])
[ {"_id":1,"num_users":56} ]
# "_id" represents postcount
```

3. Additional Ideas

Contributor statistics and gamification suggestion

The contributions of users seems incredibly skewed, possibly due to automated versus manual map editing (*the word "bot" appears in some usernames*). Here are some user percentage statistics:

- Top user contribution percentage ("jumbanho") - 52.92%
- Combined top 2 users' contribution ("jumbanho" and "woodpeck_fixbot") - 83.87%
- Combined Top 10 users contribution - 94.3%
- Combined number of users making up only 1% of posts - 287 (about 85% of all users)

Thinking about these user percentages, I'm reminded of "gamification" as a motivating force for contribution. In the context of the OpenStreetMap, if user data were more prominently displayed, perhaps others would take an initiative in submitting more edits to the map. And, if everyone sees that only a handful of power users are creating more than 90% a of given map, that might spur the creation of more efficient bots, especially if certain gamification elements were present, such as rewards, badges, or a leaderboard.

Additional data exploration using MongoDB queries

Top 10 appearing amenities

```
> db.char.aggregate([{"$match":{"amenity":{"$exists":1}}},
{"$group":{"_id":"$amenity",
"count":{"$sum":1}}}, {"$sort":{"count":1}}, {"$limit":10}])

[ {"_id":"place_of_worship","count":587},
{"_id" : "school", "count" : 416}, ....
```

Biggest religion (no surprise here)

```
> db.char.aggregate([{"$match":{"amenity":{"$exists":1},
"amenity":"place_of_worship"}},

{"$group":{"_id":"$religion", "count":{"$sum":1}}},

{"$sort":{"count":1}}, {"$limit":1}])

[ {"_id":"christian","count":577} ]
```

Most popular cuisines

```
> db.char.aggregate([{"$match":{"amenity":{"$exists":1},
"amenity":"restaurant"}}, {"$group":{"_id":"$cuisine", "count":{"$sum":1}}},
{"$sort":{"count":1}}, {"$limit":2}])

[ { "_id" : "pizza", "count" : 9},
{ "_id" : "american", "count" : 9} ...
```

Conclusion

After this review of the data it's obvious that the Charlotte area is incomplete, though I believe it has been well cleaned for the purposes of this exercise. It interests me to notice a fair amount of GPS data makes it into OpenStreetMap.org on account of users' efforts, whether by scripting a map editing bot or otherwise. With a rough GPS data processor in place and working together with a more robust data processor similar to data.pyI think it would be possible to input a great amount of cleaned data to OpenStreetMap.org.